


CISC452/CMPE452/COGS 400


Introduction to Artificial Neural Networks

Farhana Zulkernine

Learning in the Neuron



McCulloch and Pitts Neuron Model




The weights w_i take on *real values* $w_i \in \mathbb{R}$

Activation is the weighted sum of all incoming potentials.

$f(a)$ can be any function that generates a spike (high value) at a given threshold value θ to mimic the scenario of *Action Potential*.

An Artificial Neuron



The Activation Function

One possible choice is a threshold function:


Therefore, we call

this a **threshold neuron**.

$$f_i(\text{net}_i(t)) = 1, \quad \text{if } \text{net}_i(t) \geq \theta$$

$$= 0, \quad \text{otherwise}$$

1. When does it fire?
2. What output does it generate?



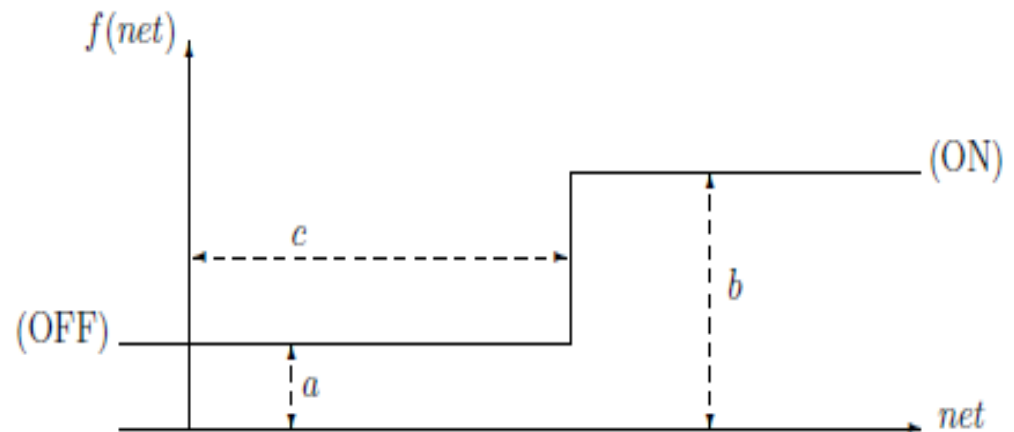
Output Functions $f(net_i)$

- The simplest node functions are:
 1. **Identity**, $f(net) = net$, and its non-negative variant $f(net) = \max(0, net)$
 2. Constant functions $f(net) = c$
 3. Signum function
$$f(net) = \begin{cases} +1 & \text{if } net > 0 \\ -1 & \text{if } net < 0 \\ 0 & \text{if } net = 0 \end{cases}$$

4. Step Function

- Simplest function that captures the idea of a "firing threshold"
- Can be used as a class identifier
- **Problem:** Very small change in $net_i(t)$ can cause a spike and hence change the output


$$\begin{aligned} f(net) &= a \text{ if } net < c \\ &= b \text{ if } net \geq c \end{aligned}$$



5. a) Ramp Function


- The ramp function is continuous and almost everywhere differentiable in exchange of the simple ON/OFF description of the output.

$$f(net) = \begin{cases} a & \text{if } net \leq c \\ a + \frac{(net-c)(b-a)}{d-c} & \text{otherwise} \\ b & \text{if } net \geq d \end{cases}$$




5. b) Piecewise Linear Functions

- Consist of finite number of linear segments, and are thus differentiable almost everywhere.
- Easier to compute than general nonlinear functions such as sigmoid functions.
- Can be used to avoid sudden change in output.



6. Sigmoid Function

- These functions are continuous and differentiable everywhere, and asymptotically approach saturation values (0 and 1 as shown in the picture)
- The parameter s controls the slope of the sigmoid function. Greater s value will give steeper curve.




$$\lim_{net \rightarrow +\infty} f(net) = 1$$

$$\lim_{net \rightarrow -\infty} f(net) = 0$$

7. Gaussian Functions

- Continuous bell-shaped functions.
- Also called 'radial-basis' function.
- $f(\text{net})$ asymptotically approaches 0 (or some constant) for large magnitudes of net, with a single maximum for $\text{net} = \mu$, say $\mu = 0$.




$$f(\text{net}) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left[-\frac{1}{2}\left(\frac{\text{net} - \mu}{\sigma}\right)^2\right]$$


Binary Analogy: Threshold Logic Units (TLUs) – AND Gate

TLUs are similar to the threshold neuron model, except that *TLUs only accept binary inputs* (0 or 1).


Example: AND Gate with $\theta = 1.5$



TLUs and Linear Separability – OR




Each point is represented by $x_1 x_2$




Input space

OR Gate: $00 = 0$ and $01, 10, 11 = 1$

TLUs and Linear Separability



Each point is represented
by $x_1 x_2$



AND Gate: 00, 01, 10 = 0 and 11 = 1

Linear Separability as Functional Mapping

- To explain linear separability, let us consider the function $f: \mathbb{R}^n \rightarrow \{0, 1\}$ with.

$$f(x_1, x_2, \dots, x_n) = \begin{cases} 1, & \text{if } \sum_{i=1}^n w_i x_i \geq \theta \\ 0, & \text{otherwise} \end{cases}$$

where x_1, x_2, \dots, x_n represent real numbers (not TLU).

This is the exactly the function that our threshold neurons use to compute their output from their inputs.


Examples – Linear Separability

Let's say we have a two-dimensional Input space (x_1 and x_2 and $n = 2$) and the following lines separate the output categories into 1 and 0.


Model the mapping using ANN.

$$f(x_1, x_2, \dots, x_n) = 1, \quad \text{if } \sum_{i=1}^n w_i x_i \geq \theta$$


$$= 0, \quad \text{otherwise}$$



For $\theta = 2$, $w_1 = 1$, $w_2 = 2$



For $\theta = 2$, $w_1 = -2$, $w_2 = 1$



For $\theta = 1$, $w_1 = -2$, $w_2 = 1$

Equation of the straight line $w_1 x_1 + w_2 x_2 = \theta$

To model it using an ANN, we need to find the weights for a given threshold.

Linear (cont...)


- Therefore, it means that by changing θ and the weights, the line can be aligned in any direction on the 2-dimensional surface to separate two categories of input based on the outputs 1 and 0.
- Here θ is called the **bias**.
- Training a network means finding the best fitting values of θ and the weights to categorize a given set of values correctly when the categories are already known.

Linear Separability (cont...)


- As we have seen, a two-dimensional input space can be divided by any straight line.
- A three-dimensional input space can be divided by any two-dimensional plane.
- In general, an *n -dimensional input space can be divided by an $(n-1)$ -dimensional plane* or hyperplane.
- Of course, for $n > 3$ this is hard to visualize.

What if the data are not linearly separable?

- A function $f: \{0, 1\}^n \rightarrow \{0, 1\}$ is linearly separable if the space of input vectors yielding 1 can be separated from those yielding 0 by a linear surface (hyperplane) in n dimensions.
- Examples (two dimensions):




linearly separable (OR)




linearly inseparable (XOR)

TLUs cannot implement XOR

What about XOR?



Each point is represented by $x_1 x_2$




XOR Gate: $00, 11=0$ and $10, 01=1$

NOT linearly separable!!!

TLUs CANNOT realize functions that are **NOT linearly separable**.

Minsky and Papert

- In 1969, Marvin Minsky and Seymour Papert, two “PSS” researchers at MIT studied the ANNs and revealed that a two-layered (input and output) network cannot handle all logical relations – specifically the XOR.
 - It implies that ANNs lacked the power of a Turing machine.
- Federal funding for ANNs immediately stopped.



Minsky


The Big Breakthrough

- David Rumelhart and Jim McClelland developed Parallel Distributed Processing (PDP)



Multilayer Networks with *Inhibition*

- The solution that Rumelhart and McClelland propose is simple: **Add a third layer between input and output.**
- 3 layers enable creating an XOR gate and handle all logic.
- Note that middle layer neuron **inhibits** output layer neuron when $x_1 = x_2 = 1$.
 - New for ANN



Activation as a Vector Product

- The net input signal is the sum of all inputs:

$$\text{net}_i(t) = \sum_{j=1}^n w_{i,j}(t) x_j(t)$$


This can be viewed as computing the **inner product** of the **vectors** \mathbf{w}_i and \mathbf{x} : https://en.wikipedia.org/wiki/Dot_product

$$\text{net}_i(t) = \| \mathbf{w}_i(t) \| \cdot \| \mathbf{x}(t) \| \cdot \cos \alpha$$

where α is the **angle** between the two vectors.

Capabilities of Threshold Neurons

- Let us assume that the threshold $\theta = 0$ and illustrate the function computed by the neuron for sample vectors w_i and x :




- Since the inner product is positive for $-90^\circ < \alpha < 90^\circ$, in this example the neuron's output is 1 for any input vector x to the right of or on the dotted line (on the same side of the line as w), and 0 for any other input vector.

What about complex functions?

- Just like Threshold Logic Units for XOR, we can combine multiple artificial neurons in multiple layers to form networks with increased capabilities.
 - For example, in a two-layer network, n input neurons in the first layer can send inputs to a single neuron in the second layer.
 - The neuron in the second layer can implement an AND function i.e., output 1 when all are 1.


Capabilities (cont...)

- What kind of function can such a network realize?



Capabilities (cont...)

- Assume that the dotted lines in the diagram represent the input-dividing lines implemented by the neurons in the first layer:




Then, for example, the second-layer neuron could output 1 if the input is within a **polygon**, and 0 otherwise.

Capabilities (cont...)

- The more neurons there are in the first layer, the more vertices can the polygons have.
- With a sufficient number of first-layer neurons, the polygons can approximate any given shape.
- The more neurons there are in the second layer, the more of these polygons can be combined to form the output function of the network.

Capabilities (cont...)


- Assume that the polygons in the diagram indicate the input regions for which each of the **second-layer** neurons yields output 1:



Then, for example, **the third-layer** neuron could output 1 if the input is within **any of the polygons**, and 0 otherwise. Example application ???

Problem


1. Draw an ANN to find if a point lies **inside** a hollow box. Explain how it would work.



2. Draw an ANN that can detect whether a point lies on the **surface** of a pyramid. Explain how it would work.

Terminology

- Example: Network function $f: \mathbb{R}^3 \rightarrow \{0, 1\}^2$ means that the network takes 3-dimensional real values as input and generates two dimensional binary output values.



Terminology

- Usually, we draw neural networks in such a way that the input enters at the bottom and the output is generated at the top.
- Arrows indicate the direction of data flow.
- The *input layer* just contains the input vector and *does not perform any computations other than distributing inputs to the next layers (used optionally)*.
- The intermediate layers, termed *hidden layers*, receives input from the input layer or previous hidden layers and *sends output to the final output layer*.
- After applying their *activation function*, the neurons in the final *output layer* generate the output vector.